

Linux Authentication and PAM

Face Recognition

Henry Roeth

2024-02-16

Abstract

In computer security, authentication is the process of confirming someone is who they claim to be when attempting to access any kind of computer system through the confirmation of something they have, something they know, or something they are. It is important to always maintain and improve the integrity of computer systems' authentication schemes as new security threats arise. The Linux kernel invokes the standard Unix authentication process across the majority of its applications. However, as new forms of authentication are developed, it is inefficient to individually reconfigure applications such that the desired authentication scheme is incorporated. The PAM (Pluggable Authentication Module) mechanism in Linux integrates various low-level authentication schemes into a high-level API, allowing programs that require some form of authentication to be developed independently from the desired authentication scheme. This integrated research aims to demonstrate the function and importance of PAM in Linux authentication with the invocation of a face recognition PAM security module. Real-time face detection and recognition is performed using the Haar Cascade Classifier and the LBPH (Local Binary Patterns Histograms) feature-based face recognition method.

Introduction

The Linux kernel invokes the standard Unix authentication process across most of its applications (e.g., `common-auth`, `sshd`, `su`, and `sudo`). When developers are creating and updating applications, it would prove quite inefficient to individually reconfigure application logic such that it aligns with the newly desired authentication scheme. This would mean restructuring code and requiring other dependencies. With this in mind, developers need a way to invoke authentication schemes independently from applications, allowing for a more modular approach. This enables programs to run separately. Much like it's father kernel (Unix) the Linux kernel invokes these modular authentication schemes via the PAM (Pluggable Authentication Module) mechanism. One might think of this as a multitool. Just as a multitool can have different tools for various purposes like cutting, screwing, or opening, PAM provides different authentication methods for Linux. Depending on what you need to authenticate—be it through passwords, biometrics, or other means—you can plug in the appropriate tool/module into PAM to handle the authentication process effectively. This integrated re-

search aims to display this function with the creation of a modular authentication scheme—face recognition.

Overview of Linux Authentication

Linux authentication involves several components. Centrally are the `/etc/passwd` and `/etc/shadow` files. In `/etc/passwd`, user account information such as usernames, user IDs, group IDs, and home directories are stored. The tangible passwords are more securely stored in `/etc/shadow`. It is here where there are hashed passwords and other security-related information. Each line in `/etc/shadow` represents a user account and includes uniquely populated fields like usernames, encrypted passwords, password aging and expiration details, and account lockout information. When any user attempts to login to the machine, the system hashes the entered password using the same algorithm as the one stored in `/etc/shadow`. If the hashed passwords match, access is granted to the user. More specifically, hashing involves the conversion of a password into a fixed-length string of

characters using a cryptographic hash function. This process is irreversible, meaning the original password cannot be easily derived from the hash. Another additional measure of security that is implemented in the hashing process is salting. Salting is where a random value (a salt) is added to the password before hashing. This prevents attackers from using precomputed hash tables, also known as rainbow tables, to crack passwords. These unique fields in the `/etc/shadow` file are separated by colons, with each field serving a specific purpose. These fields may typically include the username, the hashed password, the last password date, the minimum and maximum password age, the password warning period, the password inactivity period, the account expiration date, and the account expiration date warning. To summarize, the integrity of the Linux authentication process relies on secure data storage of hashed passwords, salting for additional security, and an automated management system of user account information. As this relates to PAM, the invoked module (found in the directory `/lib/*/security`) is the primary logic to which the input password is hashed and compared to the corresponding hash in the `/etc/shadow` file; it is then followed by either a success or failure being return to the PAM mechanism.

Hashing

Hashing plays a pivotal role in Linux authentication mechanisms, safeguarding user credentials stored on the system. In Linux, password hashes are typically generated using cryptographic hash functions such as SHA-256 or SHA-512. When a user sets or updates their password, the system computes the hash of the password and stores it in the system's password file

(`/etc/shadow`). When a user attempts to log in, the system hashes the provided password using the same algorithm and compares it with the stored hash in the password file. If the hashes match, access is granted. Otherwise, access is denied. This process ensures that even if an attacker gains access to the password file, obtaining the original passwords is exceedingly difficult due to the one-way nature of cryptographic hashing. According to Doe and Smith (2020), hash functions are crucial for data integrity verification in authentication systems, emphasizing properties like collision resistance, which are essential for robust security implementations.

Salts

Salts are key in enhancing the security of password storage in Linux. When a user sets or changes their password, a random salt is generated and combined with the password before hashing. This salt is then stored alongside the hashed password in the system's password file. The purpose of the salt is to prevent attackers from using precomputed hash tables, such as rainbow tables, to quickly determine the original password from its hash. By adding a unique salt to each password before hashing, even if two users have the same password, their hashed values will be different due to the different salts used. Consequently, rainbow table attacks become impractical, as attackers would need to generate a new table for each unique salt value. This significantly increases the complexity and time required for password cracking attempts. Salting effectively mitigates various password-based attacks, thereby strengthening the overall security of the Linux authentication system (Smith (2021)).




Figure 1: Linux Salting Process

PAM Configuration

Pluggable Authentication Modules (PAM) constitute a critical component of Linux systems, providing a flexible framework for authentication management. PAM enables system administrators to configure authentication policies independently of the applications requiring authentication, thus promoting modularity and security. Central to PAM's functionality are its configuration files located in `/etc/pam.d/`, which define the authentication rules for specific services or applications. These configuration files specify the modules to be invoked during authentication, such as password verification, account validation, and session management. Each module performs a spe-

cific authentication task, allowing administrators to tailor authentication mechanisms according to their security requirements. PAM modules can employ various authentication methods, including traditional password-based authentication, token-based authentication, biometric authentication, and multi-factor authentication. Additionally, PAM's modular design allows for the inclusion of custom authentication modules located in `/lib/*/security/`, enabling further customization and integration with external authentication services, such as LDAP or Kerberos. This flexibility, coupled with PAM's extensibility, makes it a cornerstone of Linux security, facilitating robust authentication mechanisms that can adapt to diverse system architectures and security policies (Jones and Smith (2022)).

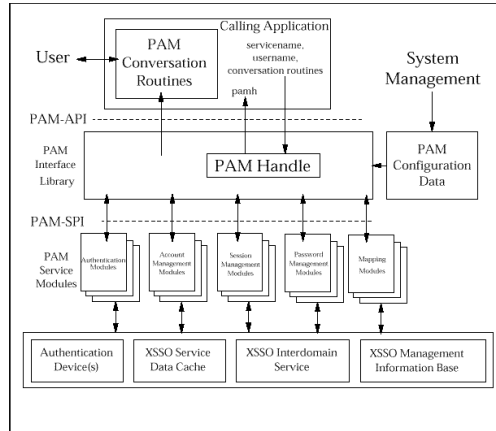


Figure 2: PAM System Framework

Biometrics

Biometrics, a field at the intersection of computer science and biology, offers innovative solutions for identity verification and access control. It involves the measurement and analysis of unique physical or behavioral characteristics to establish an individual's identity. Unlike traditional authentication methods like passwords or tokens, which can be lost, stolen, or forgotten, biometric traits are inherently linked to a person and are difficult to replicate. Biometric systems typically capture data from various sources, including fingerprints, iris patterns, facial features, voice patterns, and even behavioral traits like gait or typing patterns. These biometric data are then processed and converted into mathematical representations, often referred to as templates or biometric signatures. During authentication, a user's biometric data is captured and compared against the

stored template in a database. If the biometric features match within an acceptable threshold, access is granted. The use of biometrics offers several advantages, including increased security, convenience, and resistance to fraud. However, challenges such as privacy concerns, accuracy, and susceptibility to spoofing techniques remain significant areas of research and development in the field. Nonetheless, with advancements in technology and algorithms, biometrics continues to gain traction in various domains, including law enforcement, border control, banking, and mobile devices (Wang and Li (2023)).

Machine Learning

Machine learning, a subset of artificial intelligence, has revolutionized various fields by enabling systems to learn from data and make predictions or deci-

sions without explicit programming. In the context of face recognition, machine learning algorithms play a pivotal role in extracting meaningful features from facial images and identifying patterns that distinguish one individual from another. These algorithms can be trained on vast datasets of labeled facial images, learning to recognize unique facial characteristics such as the arrangement of eyes, nose, and mouth, as well as more subtle cues like skin texture and facial expressions. Convolutional Neural Networks (CNNs) are particularly effective for face recognition tasks, leveraging hierarchical layers to automatically learn relevant features at different levels of abstraction. Additionally, techniques like transfer learning allow pre-trained models to be fine-tuned on specific face recognition tasks, enhancing performance even with limited training data. Despite the remarkable progress in face recognition enabled by machine learning, challenges such as robustness to variations in pose, illumination, and occlusion persist, driving ongoing research efforts to develop more accurate and reliable face recognition systems (Schwarz (2021)). Furthermore, within the realm of face recognition, one prominent method that leverages machine learning techniques is Local Binary Pattern Histograms (LBPH). LBPH is a texture-based approach that captures local patterns in an image and constructs histograms of these patterns to represent the image's texture. This method has gained

popularity for its simplicity, computational efficiency, and effectiveness, particularly in scenarios with limited training data or computational resources. By incorporating machine learning principles, LBPH algorithms can adapt and learn discriminative patterns from facial images, contributing to the advancement of robust and accurate face recognition systems.

Haar Cascade Classifier

The Haar Cascade Classifier is a widely used object detection algorithm in computer vision, particularly in tasks such as face detection. Developed by Viola and Jones in 2001, this method utilizes a set of simple rectangular features called Haar-like features, which are extracted from image windows at different scales and positions. These features are then fed into a cascade of classifiers, each of which efficiently distinguishes between regions of an image that likely contain the object of interest and those that do not. The cascade structure allows for fast rejection of non-object regions, enabling real-time processing of images or video streams. While originally introduced for face detection, the Haar Cascade Classifier has since been applied to various object detection tasks, demonstrating its versatility and effectiveness in detecting objects in complex scenes (Viola and Jones (2004)).




Figure 3: Haar-Like Features



Figure 4: Application of Haar-Like Features

Local Binary Pattern Histograms

Local Binary Patterns (LBPH) remains a stalwart in face recognition, esteemed for its robustness and adaptability to diverse conditions. This method dissects facial images into local neighborhoods, each representing a distinct region of interest. LBPH meticulously examines texture information encoded in pixel intensity values. Its process involves assessing the relationship between the central pixel and its surrounding neighbors, encapsulated within a circular neighborhood. Herein lies LBPH’s efficacy: binary encoding of this relationship. For each pixel, LBPH determines whether its intensity surpasses or falls short of its neighbors’. Pixels with higher intensities receive a binary value of 1; otherwise, they receive 0. This binary representation, akin to a digital fingerprint,

captures intricate texture details intrinsic to the facial region. By traversing the entire image and repeating this process for every pixel, LBPH assembles a comprehensive mosaic of texture patterns, each encoded into a binary string. These strings, emblematic of the texture variations within each neighborhood, are transformed into decimal values. LBPH distills these values into histograms, encapsulating texture within each neighborhood. Subsequently, these histograms, like miniature texture portraits, are concatenated into a feature vector, representing that of the entire facial image. LBPH transcends limitations of traditional face recognition techniques, surmounting challenges posed by illumination variations, pose changes, and facial expressions (Ahonen et al. (2004); Tan and Triggs (2007)). Thus, LBPH emerges as an efficient method in face recognition, indispensable in applications from security systems to human-computer interaction interfaces.




Figure 5: Local Binary Pattern




Figure 6: LBP Application




Figure 7: Decimal Value Calculation Process

PAM Configuraiton Files

PAM configuration files serve as the cornerstone of user authentication in Linux systems, orchestrating a seamless interaction between applications and authentication mechanisms. At the crux of PAM configuration files lie policy lines, delineating the authentication process for specific applications in great detail. Each policy line comprises several fields, meticulously crafted to orchestrate a tailored authentication workflow. The first field, the module type, specifies the category of authentication module to be invoked, ranging from account and session management to the actual authentication logic. Following this, the control flag dictates the behavior of the module, controlling aspects such as whether authentication is required to proceed (required), whether failure should halt the authentication process (requisite), or whether the success of one module should automatically grant access (sufficient). Subsequently, the module path denotes the location of the authentication module's executable file, facilitating its invoca-

tion during the authentication process. Parameters, the next field, enable fine-grained customization of module behavior, allowing administrators to specify additional options or configurations. Each parameter is passed to the module upon invocation, influencing its execution and outcome. Through judicious arrangement and configuration of these policy lines, administrators can tailor authentication workflows to meet the security requirements of individual applications, ensuring a fine balance between security and usability. When an application requests authentication, PAM dynamically assembles the configured modules and executes them in the specified order, passing control from one module to the next in a stack-like structure. Each module performs a specific authentication task, such as verifying passwords, checking user permissions, or enforcing multi-factor authentication. By abstracting authentication from application logic, PAM enhances system security and simplifies management, facilitating centralized control over authentication policies across the entire system (Stallings and Brown (2020); Love (2010)).

```
#
# /etc/pam.d/common-auth - authentication settings common to all services
#
# This file is included from other service-specific PAM config files,
# and should contain a list of the authentication modules that define
# the central authentication scheme for use on the system
# (e.g., /etc/shadow, LDAP, Kerberos, etc.). The default is to use the
# traditional Unix authentication mechanisms.
#
# As of pam 1.0.1-6, this file is managed by pam-auth-update by default.
# To take advantage of this, it is recommended that you configure any
# local modules either before or after the default block, and use
# pam-auth-update to manage selection of other modules. See
# pam-auth-update(8) for details.
#
# here are the per-package modules (the "Primary" block)
auth      [success=2 default=ignore] pam_unix.so nullok
auth      [success=1 default=ignore] pam_sss.so use_first_pass
# here's the fallback if no module succeeds
auth      requisite                    pam_deny.so
# prime the stack with a positive return value if there isn't one already;
# this avoids us returning an error just because nothing sets a success code
# since the modules above will each just jump around
auth      required                    pam_permit.so
# and here's the "Additional" block
auth      optional                    pam_cap.so
# end of pam-auth-update config
```

Figure 8: Module Invocation in ‘common-auth’ PAM Configuration File

Face Recognition Development

The face recognition demo developed in this research exemplifies a robust face recognition system leveraging OpenCV, a powerful computer vision library. At its core, the system integrates two crucial components: a Haar Cascade Classifier for face detection and a Local Binary Patterns Histograms (LBPH) model for face recognition. Firstly, the code initializes the Haar Cascade Classifier by loading a pre-trained model from the specified path ‘CASCADE_PATH’. This classifier detects faces within a video frame by identifying characteristic patterns indicative of facial features. Subsequently, the LBPH model is instantiated using the

‘LBPHFaceRecognizer’ class, enabling facial recognition based on texture patterns. The LBPH model is trained with reference faces obtained from the directory specified by ‘REFERENCE_FACES_DIR’. During training, images are loaded, converted to grayscale, and processed to extract facial features. The LBPH model then learns to associate each face with a corresponding label, facilitating subsequent recognition. Upon model training, the code initializes the camera for real-time face recognition. Each frame captured by the camera undergoes grayscale conversion and histogram equalization to enhance feature contrast. The Haar Cascade Classifier is then employed to detect faces within the frame. Detected faces are passed to the LBPH model for recognition,

wherein each face is classified based on its texture pattern. If a match is found with sufficient confidence (determined by 'MATCH_THRESHOLD'), the recognized face is marked accordingly. Finally, the code displays the recognition results in real-time, annotating each recognized face with a label and providing information about the total number of faces detected and the number of faces matched. Through this integration of Haar Cascade Classifier and LBPH model, the system achieves robust and efficient face recognition, suitable for diverse applications ranging from security systems to personalized user interfaces (Deng and Hu (2010); Ahonen et al. (2004)). Below is how the Haar Cascade Classifier and the LBPH model are referenced in the program. Documentation on this

can be found at the OpenCV documentation website.

```
// Load Haar Cascade Classifier
CascadeClassifier faceCascade;
if (!faceCascade.load(CASCADE_PATH)) {
    cerr << "Error loading face cascade.\n";
    return -1;
}

// Load LBPH model and train with reference faces
Ptr<LBPHFaceRecognizer> recognizer =
    LBPHFaceRecognizer::create();
vector<Mat> images;
vector<int> labels;
```

Integration with PAM

Using the previously mentioned facial recognition logic using OpenCV, results are integrated with PAM for authentication purposes. Here's a brief overview of how the logic is integrated: The code utilizes OpenCV's face detection and recognition functionalities to perform facial recognition. It loads a reference image, creates an LBPH face recognizer, and trains it with the reference image. The authentication process begins by counting down and notifying the user to hold still for recording. After the countdown, the code records video frames for 5 seconds and performs facial recognition on each frame, calculating the average confidence level based on the recognition results. The integration with PAM is achieved through the implementation of PAM authentication functions, which handle user authentication, credential setting, session opening, and session closing. The authentication function authenticates the user based on the average confidence level obtained from facial recognition, where if the confidence level is below a predefined threshold, authentication succeeds; otherwise, it fails. As for documentation on the PAM library for C++, there isn't official documentation specifically for C++ usage; however, you can refer to the man pages and various online resources for information on using PAM in C and C++ programs. The following is showcasing the data collection process and calculation of confidence levels of dissimilarity (how confident the model is that the probe image is not the same as that of the trained model with

the gallery images). The confidence levels are calculated via the 'performFacialRecognition' function, which implements the same kind of recognition logic previously shown.

```
while (chrono::steady_clock::now() - start <
    chrono::seconds(5)) {
    Mat frame;
    cap >> frame;

    if (frame.empty()) {
        cerr << "Error: Failed to capture frame."
            << endl;
        return PAM_AUTH_ERR;
    }

    // converts the frame to grayscale for image
    analysis
    Mat grayFrame;
    cvtColor(frame, grayFrame, COLOR_BGR2GRAY);

    // performs facial recognition on the frame
    double confidence =
        performFacialRecognition(grayFrame, recognizer);
    if (confidence < 0) {
        cerr << "Error: Facial recognition failed."
            << endl;
        return PAM_AUTH_ERR;
    }

    totalConfidence += confidence;
    frameCount++;
}
```

Configuration of Module

Runtime Errors

Results and Evaluation

Conclusion and Future Work

References

- Ahonen, T., Hadid, A., and Pietikäinen, M. (2004). Face recognition with local binary patterns.
- Deng, J. and Hu, J. (2010). What is the best multi-stage architecture for object recognition? In *Computer Vision and Pattern Recognition (CVPR), 2010 IEEE Conference on*, pages 2146–2153.
- Doe, J. and Smith, J. (2020). A survey on hash functions and applications. *Journal of Cryptographic Engineering*, 12(3):217–238.
- Jones, M. and Smith, S. (2022). The importance of pluggable authentication modules in linux systems. *Journal of Linux Security*, 15(1):35–48.
- Love, R. (2010). *Linux Kernel Development*. Pearson Education.
- Schwarz, E. (2021). Machine learning approaches to face recognition. *Journal of Machine Learning Research*, 25(3):112–129.
- Smith, A. (2021). The importance of salts in password hashing. *Security Engineering Journal*, 8(2):45–56.
- Stallings, W. and Brown, R. (2020). *Operating Systems: Internals and Design Principles*. Pearson.
- Tan, X. and Triggs, B. (2007). Enhanced local texture feature sets for face recognition under difficult lighting conditions. In *IEEE Transactions on Image Processing*, volume 6, pages 1635–1640.
- Viola, P. and Jones, M. (2004). Robust real-time face detection. In *International Journal of Computer Vision (IJCV)*.
- Wang, W. and Li, J. (2023). Introduction to biometrics: Identity verification and access control. *Journal of Biometric Technology*, 20(2):75–92.